

HOW CAN EVERY ROLLOUT COUNT FOR SKILL SEEKERS?

COMPASS FOR SKILL PROPOSERS

Anonymous authors

Paper under double-blind review

ABSTRACT

Textual-program optimizers use execution traces to revise prompts, instructions, and agent skills. A fixed rollout budget must support both exploring new candidates and evaluating them reliably. Reusing feedback also requires tracking which examples shaped each candidate. This paper formulates *provenance-constrained evidence routing*: allocate that budget to find and select a strong program while separating the examples used to generate a revision from those used to judge it. COMPASS carries this separation through the full parent-child lineage. Its credit compares candidates on the same instances and updates whenever the highest-scoring candidates change. A persistent queue gives highly ranked candidates repeated search opportunities, and one-to-one proposal matching assigns them distinct batches with recorded room for improvement. Each fixed child is then evaluated on a separately sampled batch outside its lineage’s proposal history. Under the shared budget, more proposal attempts broaden the search, while more fresh evaluations per candidate improve selection precision. Qwen3-8B held-out evaluations span mathematical reasoning (e.g., AIME-2025), instruction following (IFBench), multi-hop question answering and fact verification (HotPotQA and HoVer), and privacy-conscious delegation (PUPA). Relative to published GEPA, the largest gain is 25.3 percentage points on AIME-2025, and the mean gain across all six shared benchmarks is 7.3 points. These results show that governing how evidence is created, reused, and routed can materially affect sample-efficient textual-program evolution.

1 INTRODUCTION

Feedback-driven optimizers revise instructions, prompts, and agent skills from execution traces. Reflection, textual gradients, and evolutionary prompt search use language feedback to produce large behavioral updates with fixed model weights (Shinn et al., 2023; Madaan et al., 2023; Yuksekgonul et al., 2025; Cheng et al., 2024; Agrawal et al., 2025). Their common optimization task is compact: given a seed textual program, a benchmark evaluator, and rollout budget B , the optimizer adaptively generates and evaluates candidates, then returns one terminal program. Every rollout serves one of two statistical roles. Proposal evidence diagnoses a failure and shapes a revision. Evaluation evidence tests a revision whose text is already fixed. Admission uses paired evaluation to decide whether the revision enters the candidate pool. Three instruction-following cases, for example, may expose a constraint-ordering error and inspire a new rule. Fresh benchmark cases then measure the fixed rule.

The number of proposal-and-evaluation attempts sets search bandwidth, and the fresh evaluations per fixed candidate determine ranking precision. Their shared budget trades broader search against more




Figure 1: **COMPASS leads on five tasks.** Qwen3-8B scores (%) use labeled center-to-rim ranges (Table 1).




Figure 2: **Legal evidence guides the next search wave in COMPASS.** Child p , proposed from parent s on P , is evaluated on a separate legal batch A and enters the cache after admission. Legal same-instance comparisons produce credit G_s , which selects proposal parents, admission references, and the final program. H matches parent slots to proposal batches, while each admission batch is sampled independently outside $X(s, P)$. Red hatching marks lineage-excluded observations.

reliable admission and final selection. Every batch that shapes a candidate becomes proposal history for its descendants. Tracking this history keeps the examples used to generate a revision separate from those used to evaluate it.

This paper studies this task as *provenance-constrained evidence routing*. The routing policy minimizes expected terminal simple regret under a total rollout budget, finite search bandwidth, finite fresh evaluation support, and lineage-level provenance. A committed partial reward table supplies legal routing evidence. The single objective is to return the strongest terminal program supported by the budget and benchmark evaluation.

Three challenges govern this objective. First, birth provenance must propagate through the full lineage so that each fixed child receives fresh admission evidence. Second, a partial and uneven reward table needs a candidate credit that remains comparable across instance difficulty and updates with new frontiers. Third, finite proposal capacity must reach distinct, high-value repair opportunities across concurrent search paths.

COMPASS (**C**redit-Aware **O**ptimization with **M**ini-batch **P**roposal and **A**dmission for **S**kill **S**earch) implements this routing principle through three aligned mechanisms. Recursive exclusion separates proposal and admission evidence throughout the lineage. Each fixed child receives a separately sampled legal admission batch. Instance-relative credit compares legally observed candidates on the same instance and refreshes their ordering whenever the reward frontier changes. A persistent queue cycles through the highest-credit candidates, including all ties at the cutoff, while a maximum-weight one-to-one allocation sends proposal slots to distinct repair opportunities. Figure 2 shows how these mechanisms route limited evidence from candidate generation to terminal selection.

Completed benchmark evaluations span model families and task types. On Qwen3-8B, task-specific COMPASS variants reach 57.33% on AIME-2025, 43.37% on IFBench, 66.67% on LiveBench-Math, and 54.00% on HoVer. These scores exceed published GEPA by 25.33, 4.76, 14.72, and 1.67 percentage points, respectively. The saved runs record the evidence and resource use behind each search decision (Sections A and D).

This paper makes four contributions:

1. *Provenance-constrained evidence routing* casts textual-program evolution as terminal simple-regret minimization under explicit constraints on rollout budget, search bandwidth, evaluation support, and lineage provenance (Section 2).
2. COMPASS recursively separates proposal from admission evidence and turns committed legal evidence into a routing signal based on how each candidate compares with its peers on the same instance. The signal is recomputed as the partial observation state changes (Sections 3.1 and 3.2).
3. A credit-ranked active set supplies persistent parent slots, and an exact maximum-weight integer assignment couples those slots to repairable proposal batches. Each fixed child receives its own legal admission batch and is compared with references on the same instances (Section 3.3).
4. COMPASS variants are evaluated under fixed published budgets across models and tasks, including mathematical reasoning, instruction following, and multi-hop fact verification (Section 5).

2 PROVENANCE-CONSTRAINED EVIDENCE ROUTING

2.1 BUDGETED TEXTUAL-PROGRAM OPTIMIZATION

Given a seed textual program p_0 , a benchmark evaluator, and rollout budget B , a routing policy π adaptively generates and evaluates candidates and returns a final program $p_{\hat{s}_\pi}$, where $\hat{s}_\pi$ is the selected candidate index. Let $R(p)$ denote expected benchmark utility and let R^* denote the best reachable utility under the search model. The budgeted objective is

$$\pi^* \in \arg \min_{\pi \in \Pi(B, b_P, b_A)} \mathbb{E}[R^* - R(p_{\hat{s}_\pi})],$$

where $b_P, b_A \in \mathbb{N}_+$, with $|P| = b_P$ proposal observations and $|A| = b_A$ fresh admission observations per attempt. For policies in $\Pi(B, b_P, b_A)$, B is the wave-level stopping threshold. A new wave starts while the cumulative proposal and evaluation costs satisfy

$$C_{\text{proposal}} + C_{\text{evaluation}} < B,$$

and runs to completion. The number $M_B(\pi)$ of completed proposal-and-evaluation attempts defines search bandwidth. Evaluation expenditure assigns b_A fresh benchmark observations to each fixed child and controls the precision of admission and terminal selection. Larger $M_B(\pi)$ lowers the search gap, and larger b_A lowers selection uncertainty.

At optimization step t , let $\mathcal{P}_t = \{p_0, \dots, p_{m_t-1}\}$ contain the accepted textual programs. Every candidate after p_0 records a parent, inducing a directed acyclic lineage. The *benchmark evaluator* runs each program using the benchmark’s official data, execution procedure, scoring metric, and aggregation rule. Executing program p_s on optimization instance x yields a bounded reward in $[0, 1]$, an output, and an execution trace. The committed cache is a partial table $q_t(s, x)$ containing the highest valid, finite score observed for each evaluated pair, with unevaluated pairs represented by missing entries. A new observation replaces an existing score and its associated output only when its score is strictly higher. Thus each candidate–instance pair retains one record. Let P_a^{birth} be the proposal batch that created candidate a , with $P_0^{\text{birth}} = \emptyset$, and let $\text{Anc}(s)$ contain the transitive ancestors of s . The lineage birth set is

$$\mathcal{B}_s = \bigcup_{a \in \text{Anc}(s) \cup \{s\}} P_a^{\text{birth}}$$

and the routing-eligible evidence set $\Omega_t(s) = \{x : q_t(s, x) \text{ exists and } x \notin \mathcal{B}_s\}$. Routing credit for candidate s therefore draws exclusively from observations outside the batches that shaped its lineage.

The benchmark evaluator supplies the optimization observations stored in the cache and, after selection, an independent frozen test.

2.2 SEPARATING PROPOSAL AND ADMISSION EVIDENCE

For a child proposed from parent s on batch P , define

$$X(s, P) = P \cup \bigcup_{a \in \text{Anc}(s) \cup \{s\}} P_a^{\text{birth}} \quad (1)$$

A legal admission batch satisfies $A \cap X(s, P) = \emptyset$. Proposal traces remain available for diagnosis. Admission for the child and its descendants uses batches disjoint from current and inherited birth provenance. In the running example, the three diagnostic cases explain the error, and a disjoint admission batch decides whether the new rule earns future search budget.

2.3 ROUTING LEGAL EVIDENCE

Let $\mathcal{H}_t$ contain the lineage, birth provenance, committed cache, remaining budget, and persistent routing state. A routing action $a_t = (\mathcal{A}_t, \mathbf{z}_t, Y_t)$ chooses active candidates, persistent parent slots, and a one-to-one proposal assignment from $\mathcal{H}_t$. The frozen benchmark test is reserved for final evaluation. Recursive exclusion separates proposal and admission evidence. Instance-relative credit organizes uneven partial evidence as frontiers change. Persistent proposal slots and a coupled one-to-one assignment coordinate concurrent repair opportunities, and each fixed child receives an independently sampled admission batch.

A feasible router preserves the recursive boundary, the wave-level stopping rule, one cached record per candidate–instance pair, and a separate final test. Its optimal policy minimizes expected simple regret for the returned program.

3 COMPASS: ROUTING LEGAL EVIDENCE THROUGH SEARCH

One COMPASS wave carries evidence through a complete lifecycle. A proposal batch reveals a failure pattern and shapes a child. Equation (1) removes the evidence that shaped its lineage from the child’s legal admission pool. A legal admission batch then judges the fixed child, eligible outcomes enter the committed cache once, and the new state determines which candidates and proposal batches receive the next budget. Provenance, evidence reuse, and allocation are linked through this shared state.

3.1 FROM PROVENANCE TO COMMITTED EVIDENCE

For parent s and proposal batch P , COMPASS samples admission instances only outside $X(s, P)$. The admission identities, references, and outcomes remain hidden from the generator until the child text is fixed. The proposal rollout may update the existing parent, but P becomes the child’s birth provenance and is therefore excluded from that child’s routing-eligible set together with all inherited birth batches. At birth, the child enters the committed table only through its legal admission evaluation. Later valid benchmark executions may enlarge its routing evidence whenever their instances remain outside the same lineage birth set. This lets proposal traces retain their high-bandwidth diagnostic role while admission supplies fresh evidence to evaluate the revision.

Let $\mathcal{F}_t(x)$ be the *reward frontier*: the candidates with the highest cached score among those for which $x \in \Omega_t(s)$, including every exact tie. Every current frontier candidate counts as a frontier outcome, irrespective of the frontier’s raw score. Define the legally observed comparison set and its frontier count by

$$\mathcal{C}_t(x) = \{u : x \in \Omega_t(u)\}, \quad N_x = |\mathcal{C}_t(x)|, \quad y_{sx} = \mathbf{1}\{s \in \mathcal{F}_t(x)\}, \quad K_x = \sum_{u \in \mathcal{C}_t(x)} y_{ux}.$$

When $N_x \geq 2$, candidate s receives the leave-one-skill-out frontier residual

$$r_{sx} = y_{sx} - \frac{K_x - y_{sx}}{N_x - 1} = \frac{N_x y_{sx} - K_x}{N_x - 1}.$$

Only comparable exposures enter the aggregate:

$$\begin{aligned} \Omega_t^{\text{cmp}}(s) &= \{x \in \Omega_t(s) : N_x \geq 2\}, \quad E_s^{\text{cmp}} = |\Omega_t^{\text{cmp}}(s)|, \\ R_s &= \sum_{x \in \Omega_t^{\text{cmp}}(s)} r_{sx}, \quad G_s = \frac{R_s}{E_s^{\text{cmp}}}. \end{aligned} \tag{2}$$

Thus G_s records how often s reaches the current frontier relative to the other legally observed candidates on the same instances. An observation with $N_x = 1$ remains in the committed table but contributes no relative credit until another candidate is legally observed on x . Missing entries remain

missing, and G_s is undefined when $E_s^{\text{cmp}} = 0$ rather than being set to zero. Whenever a commit changes an observation or a reward frontier, COMPASS recomputes N_x , K_x , all affected residuals, and the resulting G_s values.

3.2 A CREDIT-RANKED ACTIVE SET

Let $\mathcal{R}_t = \{s : E_s^{\text{cmp}} > 0\}$ be the candidates with defined credit. Parents and descendants compete in this same set. Let $\rho \in \mathbb{N}_+$ be the active-rank cutoff. If $|\mathcal{R}_t| \geq \rho$, let g_ρ be the ρ -th largest G_s and retain every candidate with $G_s \geq g_\rho$; otherwise retain all of $\mathcal{R}_t$. The active set $\mathcal{A}_t$ can therefore contain more than ρ programs at a boundary tie, and a parent and its descendant may remain active simultaneously. If no candidate has defined credit, COMPASS falls back to the seed.

The resulting active set is sorted by decreasing G_s , with earlier candidate index as the deterministic exact-tie order. Let $w \in \mathbb{N}_+$ be the wave width. A persistent cyclic cursor assigns a parent candidate to each of the w proposal slots in this order. Slot i is a distinct search opportunity, and z_i denotes its parent candidate. When $|\mathcal{A}_t| < w$, a parent can receive multiple slots in the same wave. The queue maintains coverage across waves while concentrating work on the active set.

3.3 TURNING CREDIT INTO PROPOSAL WORK AND LEGAL ADMISSION

Candidate credit selects parents for the wave’s proposal slots. The next step assigns a proposal batch to each slot. For candidate s and proposal minibatch P , define the recorded repair opportunity

$$H(s, P) = \frac{1}{|P|} \sum_{x \in P} \begin{cases} [\max_{r \in \mathcal{A}_t: q_t(r, x) \text{ finite}} q_t(r, x) - q_t(s, x)]_+, & q_t(s, x) \text{ finite,} \\ 0, & \text{otherwise.} \end{cases} \quad (3)$$

Here $[v]_+ = \max\{v, 0\}$. Missing parent observations contribute zero to this assignment weight and remain missing in the cache. Let $\mathcal{I}_w = \{0, \dots, w-1\}$ index the drawn slots z_i and proposal minibatches P_j . COMPASS allocates them by the maximum-weight binary program

$$\begin{aligned} Y^* \in \arg \max_{Y \in \{0,1\}^{w \times w}} & \sum_{i,j \in \mathcal{I}_w} H(z_i, P_j) y_{ij} \\ \text{s.t.} & \sum_{j \in \mathcal{I}_w} y_{ij} = 1 \quad \forall i, \quad \sum_{i \in \mathcal{I}_w} y_{ij} = 1 \quad \forall j. \end{aligned} \quad (4)$$

Thus distinct slots receive distinct proposal batches even when two slots carry the same parent label, and the wave realizes the largest recorded repair opportunity available to those slots and batches.

Proposal evaluation updates the existing parent’s cached observations. Tasks with usable execution traces and reflection data proceed to child generation. When perfect-score skipping is enabled, a task also requires at least one proposal score below the configured maximum. Write $m \leq w$ for the number of these remaining tasks, with parent s_i and proposal batch P_i . For each task independently, the benchmark evaluator samples an admission batch A_i satisfying

$$A_i \cap X(s_i, P_i) = \emptyset.$$

The evaluator freezes A_i and, for each $x \in A_i$, selects a reference from the current reward frontier with defined G , preferring the largest G and using seeded randomness for an exact tie. If no such comparable candidate exists, the sampled parent s_i is the reference. Denote this choice by $r_i(x)$. Missing reference observations are evaluated before the child is generated and stored using the highest-score cache rule. The admission instances, references, and outcomes remain hidden from the generator.

After the generator fixes child p_i from (s_i, P_i) , the evaluator scores it on A_i and compares its scores with the frozen reference scores. In this gate, $q(p_i, x)$ is the child’s new score and $q(r_i(x), x)$ is the bound reference score on the same instance. The paired aggregate gate is

$$\Delta_i = \sum_{x \in A_i} [q(p_i, x) - q(r_i(x), x)], \quad \text{admit}(p_i) \iff \Delta_i > 0. \quad (5)$$

Proposal and reference evaluations update existing programs’ cached scores. Only lineage-eligible entries affect their credit. An accepted child’s admission observations enter the cache with that child.

COMPASS then recomputes every affected G_s , the tie-inclusive active set, and the queue. At the budget boundary, COMPASS selects the program with the largest defined G_s , with earliest index breaking an exact tie. If every G_s is undefined, it selects the seed. The selected program then receives independent benchmark evaluation. Expanded execution and fallback details appear in Section B.

4 HOW COMPASS TURNS EVIDENCE INTO SEARCH

To see how the components fit together, follow one revision through the search loop. A proposal batch reveals a parent’s failures and shapes a child. Since those examples helped create the child, different evidence must establish that the revision is useful. The resulting fresh outcomes are compared with other candidates on the same instances, converted into routing credit, and spent on the next set of proposal opportunities. The formal statements below support three intuitions needed to read this loop: admission tests a fixed revision, credit records relative performance on shared instances, and a fixed rollout budget must trade search breadth against decision precision.

4.1 A REVISION IS JUDGED ON FRESH PAIRED EVIDENCE

A revision can look perfect on the examples it was written to solve while changing nothing elsewhere. Fresh evaluation checks whether the revision also helps beyond those examples.

Proposition 1 (Reusing proposal examples can overstate performance). *For every admission size $n \geq 1$ and $\epsilon > 0$, there is a task distribution and a proposal rule that observes n examples, constructs a program, and scores one when those examples are reused for evaluation, although the program’s population utility is at most ϵ .*

The construction in Section C memorizes the observed examples. COMPASS excludes those examples from admission for both the child that used them and later descendants that can inherit the resulting behavior.

Fresh-pair principle. The proposal batch explains how to change the parent. A disjoint admission batch decides whether that change helps. Before generation, the evaluator freezes the legal admission batch and its per-instance references and hides both from the generator. The gate in Equation (5) then compares the fixed child and its references on the same fresh instances. With n bounded pairs, the uncertainty of this paired advantage shrinks proportionally to $1/\sqrt{n}$. Thus additional admission evidence makes the decision less sensitive to a lucky minibatch. The exact non-anticipation statement and concentration bound appear in Theorem 1.

4.2 FRESH OUTCOMES BECOME INSTANCE-RELATIVE CREDIT

Admission decides whether a particular child enters the pool. Routing asks a different question: which accepted candidates have shown distinctions worth pursuing further? COMPASS answers it by comparing candidates on the same instance before aggregating across the partial cache.

Proposition 2 (Instance-wise credit conservation). *For every instance x with $N_x \geq 2$, the residuals defined immediately before Equation (2) satisfy*

$$\sum_{s \in \mathcal{C}_t(x)} r_{sx} = 0.$$

This conservation law assigns credit to differences between candidates. If all N_x compared candidates share the current frontier, the instance assigns zero residual to every candidate. If one candidate uniquely owns the frontier, it receives residual one and each of the other $N_x - 1$ candidates receives $-1/(N_x - 1)$. Thus the same frontier event means more when it distinguishes a candidate from its observed peers.

The residuals track the current comparison set. When a new committed observation changes $\mathcal{F}_t(x)$, COMPASS recomputes K_x , every affected r_{sx} , and the corresponding G_s . An instance with only one legal observation remains cached but begins contributing relative credit only after another candidate becomes comparable there.

4.3 CREDIT DIRECTS THE NEXT WAVE

Candidate credit and repair opportunity play different roles in the next wave. First, G_s determines which candidates receive the w persistent slots. Second, the recorded repair opportunity $H(s, P)$ chooses which distinct proposal batch each slot receives through the exact binary assignment in Equation (4). Credit therefore chooses *whose* search continues, while H chooses *where* each proposal attempt is spent.

Each fixed child receives an independently sampled legal batch and frozen references, and the paired gate in Equation (5) decides whether it joins the pool. Accepted evidence can change the frontiers on its instances, which redistributes the corresponding residuals and updates the next credit order. The resulting feedback loop is

$$\begin{aligned} \text{fresh outcomes} &\longrightarrow G_s \longrightarrow \text{proposal slots} \\ &\longrightarrow H\text{-matched proposals} \longrightarrow \text{fresh paired admission.} \end{aligned}$$

At the budget boundary, the program with maximal defined G_s is returned, with the seed used when every credit remains undefined. The appendix records cyclic slot coverage and exact optimality of the proposal assignment.

4.4 WHY BREADTH AND CONFIDENCE COMPETE

Let $\mathcal{P}_B$ be the candidates available when search stops at threshold B and let $\hat{p}_B$ be COMPASS’s actual G_s -selected terminal program. Its regret has the exact decomposition

$$R^* - R(\hat{p}_B) = \underbrace{R^* - \max_{p \in \mathcal{P}_B} R(p)}_{\text{search gap}} + \underbrace{\max_{p \in \mathcal{P}_B} R(p) - R(\hat{p}_B)}_{\text{selection gap}}. \quad (6)$$

The first term falls only when search discovers stronger candidates; the second falls only when the available evidence identifies a strong discovered candidate. A fixed rollout budget must support both.

A fixed-cost budget calculation makes the tradeoff explicit. Let c_{prop} be the rollout cost of producing one candidate and let $n \geq 1$ be the number of fresh evaluations used to assess it. The number of complete proposal-and-evaluation attempts that fit within B is

$$M_B(n) = \left\lfloor \frac{B}{c_{\text{prop}} + n} \right\rfloor.$$

Increasing n improves the precision of each decision but leaves room for fewer candidates. The $1/\sqrt{n}$ rate also shows diminishing returns: halving the uncertainty requires four times as many fresh evaluations.

Proposition 3 formalizes this tradeoff for a reference selector that compares candidate means under a common fresh-evaluation distribution. Its regret separates into a search gap and a selection term that shrinks with evaluation sample size. The appendix gives the full statement and assumptions. In COMPASS, credit-ranked slots and exact proposal matching allocate search effort, paired admission evaluates each revision, and the resulting G_s values guide both further search and final selection.

5 EXPERIMENTS

5.1 EVALUATION QUESTIONS

The experiments evaluate whether provenance-constrained routing improves held-out benchmark utility under a fixed budget and whether the result persists across models and tasks.

5.2 TASKS, MODELS, AND EVALUATION

The study spans AIME-2025 and LiveBench-Math mathematical reasoning, IFBench instruction following with unseen constraint types, HotPotQA question answering, HoVer multi-hop fact verification, PUPA privacy-conscious delegation (Pyatkin et al., 2025; Balunović et al., 2025; White et al., 2025; Yang et al., 2018; Jiang et al., 2020; Li et al., 2025). The models are Qwen3-8B (Yang

et al., 2025) and GPT-5.4-nano. PUPA retains its official GPT-4.1-mini judge for response quality and leakage of personally identifiable information (PII). This judge remains fixed throughout optimization and testing.

Official evaluation is separated from optimization. The Qwen3-8B AIME variant is evaluated on all 150 benchmark problems. HoVer uses 150/300/300 train, validation, and test instances with a 7,051-call budget; LiveBench-Math uses its frozen 126-problem test with a 1,839-call budget. Candidate selection uses optimization data, followed by one evaluation on the held-out test split.

5.3 QWEN3-8B: COMPARISON ACROSS SIX TASKS

Table 1 compares Qwen3-8B evaluations on all six shared tasks. The comparison rows reproduce the published GEPA study and include the initial program (Seed), Group Relative Policy Optimization (GRPO), and the prompt optimizers. The COMPASS row uses the selected program from each task-specific variant. COMPASS has the highest six-task aggregate and the highest score on five tasks. GEPA remains higher on PUPA.

Table 1: **Qwen3-8B evaluation across all shared tasks (%)**. Each cell reports the official task aggregate. Aggregate is the unweighted mean over the six tasks, and Δ Seed is the aggregate improvement in percentage points. Comparison rows are published GEPA-study results; TextGrad and Trace were not reported for Qwen3-8B.

Method	HotPotQA	IFBench	HoVer	PUPA	AIME-2025	LiveBench-Math	Aggregate	Δ Seed
Seed	42.33	36.90	35.33	80.82	27.33	48.70	45.23	–
GRPO	43.33	35.88	38.67	86.66	38.00	51.26	48.91	+3.68
MIPROv2	55.33	36.22	47.33	81.55	20.00	46.60	47.84	+2.61
GEPA	62.33	38.61	52.33	91.85	32.00	51.95	54.85	+9.62
GEPA+Merge	64.33	28.23	51.67	86.26	32.00	51.95	52.40	+7.17
COMPASS	64.67	43.37	54.00	86.91	57.33	66.67	62.16	+16.93

The six tests contain 300, 294, 300, 221, 150, and 126 instances for HotPotQA, IFBench, HoVer, PUPA, AIME-2025, and LiveBench-Math, respectively. Candidate selection precedes the official test. The selected-candidate records are listed in Section D.

5.4 GPT-5.4-NANO: SAME-MODEL COMPARISON

Table 2 reports the GPT-5.4-nano comparison. Both task execution and reflection use GPT-5.4-nano. PUPA keeps the fixed benchmark scorer described above. COMPASS has the highest six-task aggregate and the highest score on four tasks. MIPROv2 is higher on HotPotQA, and GEPA is higher on HoVer.

Table 2: **GPT-5.4-nano evaluation across all shared tasks (%)**. Each cell reports the official task aggregate. Aggregate is the unweighted mean over the six tasks, and Δ Seed is the aggregate improvement in percentage points. The table includes methods with completed same-model evaluations.

Method	HotPotQA	IFBench	HoVer	PUPA	AIME-2025	LiveBench-Math	Aggregate	Δ Seed
Seed	31.33	48.98	45.33	74.86	15.33	35.09	41.82	–
MIPROv2	61.67	47.79	53.00	75.45	30.00	43.66	51.93	+10.11
GEPA	57.33	48.30	61.33	81.09	34.00	35.03	52.85	+11.03
COMPASS	60.00	51.36	60.33	84.14	42.00	45.03	57.14	+15.32

6 RELATED WORK

Reflective textual-program optimization. Reflexion and Self-Refine use language feedback to revise behavior, while TextGrad, Trace/OptoPrime, MIPROv2, and GEPA optimize textual components from execution evidence (Shinn et al., 2023; Madaan et al., 2023; Yuksekogunul et al., 2025;

Cheng et al., 2024; Opsahl-Ong et al., 2024; Agrawal et al., 2025). COMPASS focuses on evidence routing: which observations are eligible to evaluate each candidate, how its frontier contribution should be measured against other legally observed candidates on the same instance, and where the next proposal budget should flow.

Evolution, illumination, and skill search. Evolutionary prompt and skill methods maintain multiple candidates to avoid a single greedy trajectory (Agrawal et al., 2025; Mi et al., 2026; Moll et al., 2026; Zhang et al., 2026; Wang et al., 2026b). COMPASS keeps parents and descendants in one credit-ranked candidate pool. Its active set contains the highest-credit candidates, including all ties at the cutoff.

Adaptive evaluation and budget allocation. TRIPLE, POES, and Bayesian-Agent allocate evaluation or retrieval effort under limited budgets (Shi et al., 2024; Ma et al., 2026; Wu et al., 2026); TextBO and APEX optimize evaluation efficiency and development-example choice (Kang & Yoganarasimhan, 2026; Wang et al., 2026a). COMPASS couples instance-relative candidate credit, exact one-to-one proposal assignment, and independently sampled legal admission in one shared state. Its credit measures each candidate’s frontier membership relative to the other candidates observed on the same instance.

Generalization after adaptive search. Repeated adaptive use of optimization data can widen the gap between selected validation performance and fresh test performance (Dwork et al., 2015). GEPA addresses this empirically through reusable rules and branch diversity (Agrawal et al., 2025). COMPASS adds recursive separation of proposal and admission evidence and an instance-relative frontier residual that is recomputed as the cached observations change. Held-out benchmark tests evaluate the selected program.

7 CONCLUSION

This paper treats textual-program evolution as a fixed-budget information problem. Fresh evaluations reduce decision noise, new candidates broaden the search, and provenance determines which evidence can evaluate a lineage. COMPASS links these decisions through recursive exclusion, instance-relative credit, a credit-ranked active set, exact proposal assignment, and legal admission. Qwen3-8B and GPT-5.4-nano evaluations cover all six shared tasks. The shared routing state determines which observations contribute credit and where that credit directs the next search effort.

REFERENCES

- Lakshya A. Agrawal, Shangyin Tan, Dilara Soylu, Noah Ziemis, Rishi Khare, Krista Opsahl-Ong, Arnav Singhvi, Herumb Shandilya, Michael J. Ryan, Meng Jiang, Christopher Potts, Koushik Sen, Alexandros G. Dimakis, Ion Stoica, Dan Klein, Matei Zaharia, and Omar Khattab. GEPA: Reflective prompt evolution can outperform reinforcement learning, 2025. URL <https://arxiv.org/abs/2507.19457>. ICLR 2026 Oral.
- Mislav Balunović, Jasper Dekoninck, Ivo Petrov, Nikola Jovanović, and Martin Vechev. Aime-2025, February 2025. URL https://huggingface.co/datasets/MathArena/aime_2025. Dataset from MathArena: Evaluating LLMs on Uncontaminated Math Competitions.
- Ching-An Cheng, Allen Nie, and Adith Swaminathan. Trace is the next autodiff: Generative optimization with rich feedback, execution traces, and llms, 2024. URL <https://arxiv.org/abs/2406.16218>.
- Cynthia Dwork, Vitaly Feldman, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Aaron Roth. Generalization in adaptive data analysis and holdout reuse. In *Advances in Neural Information Processing Systems*, volume 28, 2015. URL <https://proceedings.neurips.cc/paper/2015/hash/bad5f33780c42f2588878a9d07405083-Abstract.html>.
- Yichen Jiang, Shikha Bordia, Zheng Zhong, Charles Dognin, Maneesh Singh, and Mohit Bansal. HoVer: A dataset for many-hop fact extraction and claim verification. In Trevor Cohn, Yulan He, and Yang Liu (eds.), *Findings of the Association for Computational Linguistics: EMNLP*

2020, pp. 3441–3460, Online, November 2020. Association for Computational Linguistics. doi: 10.18653/v1/2020.findings-emnlp.309. URL <https://aclanthology.org/2020.findings-emnlp.309/>.

Enoch Hyunwook Kang and Hema Yoganasimhan. TextBO: Bayesian optimization in language space for eval-efficient self-improving AI. In *International Conference on Learning Representations*, 2026. URL <https://arxiv.org/abs/2511.12063>.

Siyan Li, Vethavikashini Chithrra Raghuram, Omar Khattab, Julia Hirschberg, and Zhou Yu. PAILLON: Privacy preservation from Internet-based and local language model ensembles. In Luis Chiruzzo, Alan Ritter, and Lu Wang (eds.), *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pp. 3371–3390, Albuquerque, New Mexico, April 2025. Association for Computational Linguistics. ISBN 979-8-89176-189-6. doi: 10.18653/v1/2025.naacl-long.173. URL <https://aclanthology.org/2025.naacl-long.173/>.

Xiaoyu Ma, Yiwen Li, Haoyue Liu, Zhichao Wang, Ye Chen, Yongxin Guo, and Xiaoying Tang. Select smarter, not more: Prompt-aware evaluation scheduling with submodular guarantees, 2026. URL <https://arxiv.org/abs/2604.11328>.

Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wier-effe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Shashank Gupta, Bodhisattwa Prasad Majumder, Katherine Hermann, Sean Welleck, Amir Yazdanbakhsh, and Peter Clark. Self-refine: Iterative refinement with self-feedback. In A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine (eds.), *Advances in Neural Information Processing Systems*, volume 36, pp. 46534–46594. Curran Associates, Inc., 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/file/91eddf07232fblb55a505a9e9f6c0ff3-Paper-Conference.pdf.

Qirui Mi, Zhijian Ma, Mengyue Yang, Haoxuan Li, Yisen Wang, Haifeng Zhang, and Jun Wang. Skill-Pro: Learning reusable skills from experience via non-parametric PPO for LLM agents, 2026. URL <https://arxiv.org/abs/2602.01869>. ICML 2026 Spotlight.

Johannes Moll, Jean-Philippe Corbeil, Jiazhen Pan, Martin Hadamitzky, Daniel Rueckert, Lisa Adams, and Keno Bressem. GRASP: Gated regression-aware skill proposer for self-improving LLM agents, 2026. URL <https://arxiv.org/abs/2605.29668>.

Krista Opsahl-Ong, Michael J Ryan, Josh Purtell, David Broman, Christopher Potts, Matei Zaharia, and Omar Khattab. Optimizing instructions and demonstrations for multi-stage language model programs. In Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen (eds.), *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 9340–9366, Miami, Florida, USA, November 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.emnlp-main.525. URL <https://aclanthology.org/2024.emnlp-main.525/>.

Valentina Pyatkin, Saumya Malik, Victoria Graf, Hamish Ivison, Shengyi Huang, Pradeep Dasigi, Nathan Lambert, and Hannaneh Hajishirzi. Generalizing verifiable instruction following, 2025. URL <https://arxiv.org/abs/2507.02833>.

Chengshuai Shi, Kun Yang, Zihan Chen, Jundong Li, Jing Yang, and Cong Shen. Efficient prompt optimization through the lens of best arm identification. In *Advances in Neural Information Processing Systems*, 2024. URL <https://openreview.net/forum?id=FLNnlFBGMo>.

Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: language agents with verbal reinforcement learning. In *Proceedings of the 37th International Conference on Neural Information Processing Systems, NIPS ’23*, Red Hook, NY, USA, 2023. Curran Associates Inc.

Fei Wang, Si Si, Cho-Jui Hsieh, and Inderjit S. Dhillon. APEX: Automated prompt engineering expert with dynamic data selection, 2026a. URL <https://arxiv.org/abs/2606.11459>.

Zefeng Wang, Minxi Yan, Jinhe Bi, Sikuan Yan, Volker Tresp, and Yunpu Ma. MetaSkill-Evolve: Recursive self-improvement of LLM agents via two-timescale meta-skill evolution, 2026b. URL <https://arxiv.org/abs/2607.05297>.

-
- Colin White, Samuel Dooley, Manley Roberts, Arka Pal, Ben Feuer, Siddhartha Jain, Ravid Shwartz-Ziv, Neel Jain, Khalid Saifullah, Sreemanti Dey, Shubh-Agrawal, Sandeep Singh Sandha, Siddhartha Naidu, Chinmay Hegde, Yann LeCun, Tom Goldstein, Willie Neiswanger, and Micah Goldblum. Livebench: A challenging, contamination-limited llm benchmark, 2025. URL <https://arxiv.org/abs/2406.19314>.
- Xiaojun Wu, Cehao Yang, Honghao Liu, Xueyuan Lin, Wenjie Zhang, Zhichao Shi, Xuhui Jiang, Chengjin Xu, Jia Li, and Jian Guo. Bayesian-Agent: Posterior-guided skill evolution for LLM agent harnesses, 2026. URL <https://arxiv.org/abs/2606.08348>.
- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Haoran Wei, Huan Lin, Jialong Tang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiayi Yang, Jing Zhou, Jingren Zhou, Junyang Lin, Kai Dang, Keqin Bao, Kexin Yang, Le Yu, Lianghao Deng, Mei Li, Mingfeng Xue, Mingze Li, Pei Zhang, Peng Wang, Qin Zhu, Rui Men, Ruize Gao, Shixuan Liu, Shuang Luo, Tianhao Li, Tianyi Tang, Wenbiao Yin, Xingzhang Ren, Xinyu Wang, Xinyu Zhang, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yinger Zhang, Yu Wan, Yuqiong Liu, Zekun Wang, Zeyu Cui, Zhenru Zhang, Zhipeng Zhou, and Zihan Qiu. Qwen3 technical report, 2025. URL <https://arxiv.org/abs/2505.09388>.
- Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2018.
- Mert Yuksekgonul, Federico Bianchi, Joseph Boen, Sheng Liu, Pan Lu, Zhi Huang, Carlos Guestrin, and James Zou. Optimizing generative ai by backpropagating language model feedback. *Nature*, 639:609–616, 2025.
- Hanrong Zhang, Shicheng Fan, Henry Peng Zou, Yankai Chen, Zhenting Wang, Jiayu Zhou, Chengze Li, Wei-Chieh Huang, Yifei Yao, Kening Zheng, Xue Liu, Xiaoxiao Li, and Philip S. Yu. CoEvoSkills: Self-evolving agent skills via co-evolutionary verification, 2026. URL <https://arxiv.org/abs/2604.01687>.

A NOTATION AND EXACT STATE

A.1 ELIGIBLE EXPOSURE AND FRONTIERS

At a fixed routing state t , reuse the lineage-eligible set $\Omega_t(s) = \{x : q_t(s, x) \text{ exists and } x \notin \mathcal{B}_s\}$ from Section 2. The reward frontier is

$$\mathcal{F}_t(x) = \arg \max_{s: x \in \Omega_t(s)} q_t(s, x),$$

with all exact ties retained and every frontier candidate counted, regardless of the frontier’s raw score. For each instance, define

$$\mathcal{C}_t(x) = \{u : x \in \Omega_t(u)\}, \quad N_x = |\mathcal{C}_t(x)|, \quad y_{sx} = \mathbf{1}\{s \in \mathcal{F}_t(x)\}, \quad K_x = \sum_{u \in \mathcal{C}_t(x)} y_{ux}.$$

When $N_x \geq 2$, the per-observation credit is

$$r_{sx} = y_{sx} - \frac{K_x - y_{sx}}{N_x - 1}.$$

The candidate aggregate is

$$\Omega_t^{\text{cmp}}(s) = \{x \in \Omega_t(s) : N_x \geq 2\}, \quad E_s^{\text{cmp}} = |\Omega_t^{\text{cmp}}(s)|, \quad R_s = \sum_{x \in \Omega_t^{\text{cmp}}(s)} r_{sx}, \quad G_s = R_s / E_s^{\text{cmp}}.$$

An $N_x = 1$ observation remains committed but does not enter the aggregate. Whenever a new finite observation changes $\mathcal{C}_t(x)$ or $\mathcal{F}_t(x)$, the router recomputes N_x , K_x , and every affected candidate contribution on that instance. Candidates with $E_s^{\text{cmp}} = 0$ have undefined credit; if every credit is undefined, routing and final selection fall back to the seed.

A.2 RECURSIVE PROPOSAL-LINEAGE EXCLUSION

If parent s is proposed on batch P , the prospective child cannot use its current proposal instances or any ancestor birth-proposal instances as admission evidence. Let P_a^{birth} be the recorded proposal batch that created accepted candidate a , with $P_0^{\text{birth}} = \emptyset$ for the seed. The excluded set is

$$X(s, P) = P \cup \bigcup_{a \in \text{Anc}(s) \cup \{s\}} P_a^{\text{birth}}.$$

Birth provenance records the proposal batches that shaped a candidate and its ancestors. Admission observations retain their role as evaluation evidence. The lineage birth set $\mathcal{B}_s$ governs which committed observations may contribute to candidate s ’s routing key; $X(s, P)$ extends that exclusion to the prospective child’s admission.

A.3 CHECKPOINT STATE

The optimizer checkpoint records the accepted programs, parent relation, committed reward cache, frontier map, logical call count, active-queue cursor, proposal assignments, sampled admission batches, and seeded random states. The queue cursor is interpreted modulo the current active order after every recomputation.

B ALGORITHMIC DETAILS

B.1 PROPOSAL ALLOCATION IMPLEMENTATION

Equations (3) and (4) define the proposal objective. The implementation builds the complete slot–batch weight matrix from the committed cache and solves the resulting maximum-weight assignment exactly. Duplicate parent labels remain distinct left-side slot vertices. Missing cache entries remain unchanged during assignment.

B.2 DIRECT LEGAL ADMISSION IMPLEMENTATION

For each of the $m \leq w$ remaining proposal tasks, the benchmark evaluator samples one batch A_i from that task’s admission sampler after removing every instance in $X(s_i, P_i)$. For each instance in the sampled batch, the evaluator selects the frontier candidate with defined, maximal G as reference, using seeded randomness for an exact tie and the sampled parent when no comparable candidate exists. The resulting plan is hidden from the generator, required missing reference observations are filled on the existing reference programs, and the fixed child is admitted exactly when the paired aggregate in Equation (5) is positive.

B.3 WAVE EXECUTION

Algorithm 1 One COMPASS wave

Require: committed routing state $\mathcal{H}_t$, wave width w , active cutoff ρ
Ensure: updated committed state $\mathcal{H}_{t+1}$

- 1: Recompute every defined G_s from the committed cache
- 2: $\mathcal{A}_t \leftarrow$ tie-inclusive Top- ρ candidates, sorted by G_s and candidate index
- 3: If $\mathcal{A}_t = \emptyset$, set $\mathcal{A}_t \leftarrow \{0\}$ ▷ Seed fallback
- 4: $(z_i)_{i \in \mathcal{I}_w} \leftarrow \text{CYCLICSLOTS}(\mathcal{A}_t, w)$
- 5: Draw $(P_j)_{j \in \mathcal{I}_w}$ and form $M_{ij} \leftarrow H(z_i, P_j)$
- 6: $Y^* \leftarrow \text{MAXWEIGHTASSIGNMENT}(M)$
- 7: **for all** assigned parent-proposal pairs (s_i, P_i) **in parallel do**
- 8: Evaluate s_i on P_i and construct its reflection data
- 9: Update the cache with new pairs or strictly higher parent scores
- 10: **end for**
- 11: $\mathcal{T} \leftarrow$ tasks with usable traces and reflection data, after configured perfect-score skipping
- 12: $m \leftarrow |\mathcal{T}| \leq w$
- 13: **for all** $(s_i, P_i) \in \mathcal{T}$ **in parallel do**
- 14: $A_i \leftarrow \text{SAMPLEADMISSION}(s_i, P_i)$ subject to $A_i \cap X(s_i, P_i) = \emptyset$
- 15: $r_i \leftarrow \text{BINDREFERENCES}(A_i)$; freeze and hide (A_i, r_i) from the generator
- 16: Evaluate missing references and store their scores and outputs
- 17: $p_i \leftarrow \text{REFLECT}(s_i, P_i)$; evaluate p_i on A_i
- 18: $\Delta_i \leftarrow \sum_{x \in A_i} [q(p_i, x) - q(r_i(x), x)]$
- 19: **if** $\Delta_i > 0$ **then**
- 20: Commit p_i and its lineage-eligible admission evidence
- 21: **end if**
- 22: **end for**
- 23: Update the logical budget, frontiers, credits, and queue cursor; persist $\mathcal{H}_{t+1}$

Repeat Algorithm 1 while the cumulative logical call count is below B , and complete the final wave before stopping. Report both the target B and actual calls. Request-level retries add no logical task result.

C PROOFS FOR EVIDENCE, CREDIT, AND REGRET

The results below formalize three parts of the mechanism: recursive separation provides fresh paired admission evidence, relative credit is conserved within each comparable instance, and selection by fresh sample means demonstrates the budget tradeoff between search breadth and decision precision.

C.1 WHY PROPOSAL AND ADMISSION EVIDENCE ARE SEPARATE

Proof of Proposition 1. Let the task distribution be uniform on a domain of size $K \geq n/\epsilon$. After observing $S = (x_1, \dots, x_n)$, let the proposal rule construct a program whose reward is one exactly on the distinct elements appearing in S and zero elsewhere. Every reused evaluation example belongs to that set, so the empirical reward on S is one. The population reward is the fraction of the domain covered by the observed set, at most $n/K \leq \epsilon$. $\square$

The construction motivates admission on examples outside those that shaped the candidate. Recursive birth provenance extends this separation through the lineage, because a descendant can retain behavior adapted to an ancestor's proposal batch.

C.2 RECURSIVE EXCLUSION GIVES FRESH PAIRED ADMISSION

Theorem 1 (Fresh paired admission). *Let $\mathcal{V}$ be the information visible to the generator when child p is fixed from parent s and proposal batch P , and suppose that p is $\mathcal{V}$ -measurable. Before generation, the evaluator fixes an admission batch $A = (X_1, \dots, X_n)$ outside $X(s, P)$ and binds each reference $r(X_j)$ from the committed pre-child state and seeded tie-breaking randomness. Admission identities, references, and outcomes are not visible in $\mathcal{V}$. Then*

$$A \cap P = \emptyset, \quad A \cap P_a^{\text{birth}} = \emptyset \quad \text{for every } a \in \text{Anc}(s) \cup \{s\}.$$

After p is fixed, reveal the paired outcomes sequentially. Let $\mathcal{H}_{j-1}$ contain the fixed child, the evaluator's binding state, and the first $j-1$ revealed pairs, and define

$$Z_j = q(p, X_j) - q(r(X_j), X_j), \quad \mu_j = \mathbb{E}[Z_j \mid \mathcal{H}_{j-1}],$$

$$\hat{\mu}_A(p; r) = \frac{1}{n} \sum_{j=1}^n Z_j, \quad \bar{\mu}_A(p; r) = \frac{1}{n} \sum_{j=1}^n \mu_j.$$

Because $Z_j \in [-1, 1]$, for every $\delta \in (0, 1)$, with probability at least $1 - \delta$,

$$|\hat{\mu}_A(p; r) - \bar{\mu}_A(p; r)| \leq \sqrt{\frac{2 \log(2/\delta)}{n}}. \quad (7)$$

For the i th task, the implemented gate is equivalently admit(p_i) $\iff \hat{\mu}_{A_i}(p_i; r_i) > 0$.

Proof of Theorem 1. Because $X(s, P)$ contains both P and every P_a^{birth} for $a \in \text{Anc}(s) \cup \{s\}$, legality $A \cap X(s, P) = \emptyset$ gives the displayed disjointness relations. After the child is fixed, each paired difference Z_j lies in an interval of length two. Conditional Hoeffding–Azuma applied to $Z_j - \mathbb{E}[Z_j \mid \mathcal{H}_{j-1}]$ gives

$$\Pr(|\hat{\mu}_A(p; r) - \bar{\mu}_A(p; r)| \geq \epsilon) \leq 2 \exp\left(-\frac{n\epsilon^2}{2}\right).$$

Setting the right-hand side to δ proves Equation (7). $\square$

Equation (1) supplies this separation through the entire birth lineage. COMPASS binds admission before child generation and hides admission identities and references from the generator. The resulting paired contrast supplies conditionally fresh evidence for the admission gate.

C.3 RELATIVE CREDIT IS CONSERVED INSTANCE-WISE

Proof of Proposition 2. For a fixed comparable instance x , use $\sum_{s \in C_t(x)} y_{sx} = K_x$ to obtain

$$\begin{aligned} \sum_{s \in C_t(x)} r_{sx} &= K_x - \frac{\sum_{s \in C_t(x)} (K_x - y_{sx})}{N_x - 1} \\ &= K_x - \frac{N_x K_x - K_x}{N_x - 1} = 0. \end{aligned}$$

The statement is per instance; the candidate averages G_s need not sum to zero because their comparable exposure sets and denominators can differ. $\square$

C.4 THE SEARCH–EVALUATION BUDGET TRADEOFF

The reference calculation in Section 4.4 compares candidates using the mean of fresh evaluation scores. The following statement gives its selection error and separates that error from the search gap.

Proposition 3 (Fresh-evidence breadth–confidence tradeoff). *For any $M, n \in \mathbb{N}_+$ and $\delta \in (0, 1)$, let $p_1, \dots, p_M$ be adaptively generated candidates. After p_i is fixed, suppose it receives n conditionally independent and identically distributed benchmark rewards in $[0, 1]$ under one common evaluation distribution, with mean $R(p_i)$. Let $\hat{R}_n(p_i)$ be their mean and define*

$$\tilde{p}_{M,n} \in \arg \max_{p \in \{p_1, \dots, p_M\}} \hat{R}_n(p).$$

Then, with probability at least $1 - \delta$,

$$R^* - R(\tilde{p}_{M,n}) \leq \underbrace{R^* - \max_{i \leq M} R(p_i)}_{\text{search gap}} + 2 \underbrace{\sqrt{\frac{\log(2M/\delta)}{2n}}}_{\text{selection noise}}. \quad (8)$$

For Proposition 3, let $\mathcal{F}_{i-1}$ contain the complete adaptive search history before candidate p_i is evaluated and let p_i be measurable with respect to that history. Conditional on $\mathcal{F}_{i-1}$, its n evaluation rewards follow the distribution in the proposition and are hidden until p_i is fixed. The reference selector $\tilde{p}_{M,n}$ chooses a program with the largest fresh sample mean.

Proof of Proposition 3. Conditional on $\mathcal{F}_{i-1}$ and p_i , Hoeffding’s inequality bounds candidate i ’s empirical error by ϵ with failure probability $2e^{-2n\epsilon^2}$. Taking expectations over the history preserves this bound, and a union bound over M candidates yields simultaneous error at most $\epsilon = \sqrt{\log(2M/\delta)/(2n)}$. Let p_M^* maximize population utility among the generated candidates. On this event, $R(p_M^*) - R(\tilde{p}_{M,n}) \leq 2\epsilon$. Adding $R^* - R(p_M^*)$ gives Equation (8). $\square$

Because rewards lie in $[0, 1]$, taking expectations and bounding the complement of the concentration event by one also gives

$$\mathbb{E}[R^* - R(\tilde{p}_{M,n})] \leq \mathbb{E}\left[R^* - \max_{i \leq M} R(p_i)\right] + 2\sqrt{\frac{\log(2M/\delta)}{2n}} + \delta.$$

The two terms isolate the reference tradeoff. Increasing n tightens each evaluation decision but lowers $M_B(n)$; increasing bandwidth raises the chance of discovering a strong program but leaves less fresh support per candidate. Along a growing candidate sequence, the expected search-regret term is nonincreasing in M , whereas the displayed selection term decreases with n . In COMPASS, recursive exclusion, instance-relative candidate ranking, cyclic slots, and H -based proposal matching route the corresponding budget across fresh decisions and new proposal attempts.

C.5 STRUCTURAL PROPERTIES OF THE ROUTER

Proposition 4 (Cyclic slot coverage). *For a fixed ordered active set of size m , every block of m consecutive queue draws contains each active candidate exactly once, including when a width- w wave wraps around the queue.*

Proof of cyclic slot coverage. Index the fixed active order by $0, \dots, m-1$ and let the persistent cursor be c . Consecutive draws are $(c+r) \bmod m$. For $r = 0, \dots, m-1$, these residues are pairwise distinct and cover every index. Grouping the same sequence into blocks of width w does not alter the argument, including a block that wraps around the end of the queue. $\square$

Proposition 5 (Wave-optimal proposal matching). *Conditional on the w drawn parent slots and w proposal minibatches, Equation (4) maximizes total recorded repairable gap over all one-to-one slot-minibatch assignments.*

Proof of wave-optimal proposal matching. The w parent slots and w minibatches define a complete bipartite graph with edge weight $H(z_i, P_j)$. A feasible binary matrix in Equation (4) is exactly a perfect matching, or equivalently a permutation matrix. Selecting the maximum-weight feasible matrix therefore returns the maximum-weight matching. Repeated parent labels do not change the graph because their slot vertices remain distinct. $\square$

D DETAILED EXPERIMENTAL PROTOCOL

D.1 METHOD CONFIGURATION

The method in Section 3 uses instance-relative frontier credit, applies the tie-inclusive Top- ρ cutoff to all candidates with defined credit, uses the exact proposal assignment in Equation (4), and samples admission batches independently for each remaining proposal task. The implementation names this credit `leave_one_out_frontier_credit`. If the final selector returns seed index zero, the protocol records that the seed was retained and skips a new official evaluation. Otherwise, it evaluates the selected program once.

The Qwen3-8B runs select candidates 55, 78, 76, 34, 15, and 30 for HotPotQA, IFBench, HoVer, PUPA, AIME-2025, and LiveBench-Math, respectively. GEPA’s GPT-5.4-nano AIME-2025 result uses its first completed official evaluation. Each saved run records rollout counts, candidate indices, tokens, and costs. Its checkpoint retains lineage, queue, assignment, and cache updates for analysis of the search mechanism.

D.2 OFFICIAL EVALUATION PROTOCOL

Reported results use the candidate fixed at the end of optimization and a complete official evaluation. Each run records the selected program, source version, configuration, and model. IFBench requires 294 finite example scores whose mean equals the aggregate. AIME requires all 150 benchmark problems. HoVer requires 300 finite binary scores, and LiveBench-Math requires all 126 frozen test positions under its official aggregation.

Evaluation uses the selected candidate with its matching implementation and configuration. Service failures are recorded separately from benchmark scores. Preflight runs and intermediate checkpoints remain separate from the final evaluation, and the candidate is fixed before its test scores are observed.

D.3 RESOURCE ACCOUNTING

Logical task calls and final-wave overshoot are reported separately from physical language-model calls. The latter include input/output tokens and route-priced cost. A local task deadline can produce a benchmark failure score when that outcome is defined by the evaluator. Other provider, schema, serialization, and engineering errors terminate the evaluation batch. A resumable run must preserve the exact source, config, cache, checkpoint, and random-number-generator state.

E BROADER IMPACT

Evidence routing can reduce the number of expensive environment executions needed to discover a useful textual program and can make search decisions easier to audit. The same efficiency can also optimize undesirable procedures or exploit imperfect rewards. Responsible deployment calls for secure tool policies, sandboxing, red-team evaluation, and human review.

Use of large language models. Language models assisted with literature review, editing, code, and figure design. The authors verified every citation, algorithm, and artifact; model output was neither experimental evidence nor used to choose test candidates.